Channel Configuration (configtx)

Shared configuration for a Hyperledger Fabric blockchain network is stored in a collection configuration transactions, one per channel. Each configuration transaction is usually referred to by the shorter name *configtx*.

Channel configuration has the following important properties:

- 1. **Versioned**: All elements of the configuration have an associated version which is advanced with every modification. Further, every committed configuration receives a sequence number.
- 2. **Permissioned**: Each element of the configuration has an associated policy which governs whether or not modification to that element is permitted. Anyone with a copy of the previous configtx (and no additional info) may verify the validity of a new config based on these policies.
- 3. **Hierarchical**: A root configuration group contains sub-groups, and each group of the hierarchy has associated values and policies. These policies can take advantage of the hierarchy to derive policies at one level from policies of lower levels.

Anatomy of a configuration

Configuration is stored as a transaction of type HeaderType_CONFIG in a block with no other transactions. These blocks are referred to as Configuration Blocks, the first of which is referred to as the Genesis Block.

```
The proto structures for configuration are stored in <a href="fabric/protos/common/configtx.proto">fabric/protos/common/configtx.proto</a>.

The Envelope of type <a href="HeaderType_CONFIG">HeaderType_CONFIG</a> encodes a <a href="ConfigEnvelope">ConfigEnvelope</a> is defined as follows:
```

```
message ConfigEnvelope {
   Config config = 1;
   Envelope last_update = 2;
}
```

The last_update field is defined below in the **Updates to configuration** section, but is only necessary when validating the configuration, not reading it. Instead, the currently committed configuration is stored in the config field, containing a config message.

```
message Config {
    uint64 sequence = 1;
    ConfigGroup channel_group = 2;
}
```

The sequence number is incremented by one for each committed configuration. The channel_group field is the root group which contains the configuration. The ConfigGroup structure is recursively defined, and builds a tree of groups, each of which contains values and policies. It is defined as follows:

```
message ConfigGroup {
    uint64 version = 1;
    map<string,ConfigGroup> groups = 2;
    map<string,ConfigValue> values = 3;
    map<string,ConfigPolicy> policies = 4;
    string mod_policy = 5;
}
```

Because ConfigGroup is a recursive structure, it has hierarchical arrangement. The following example is expressed for clarity in golang notation.

```
// Assume the following groups are defined
var root, child1, child2, grandChild1, grandChild2, grandChild3 *ConfigGroup
// Set the following values
root.Groups["child1"] = child1
root.Groups["child2"] = child2
child1.Groups["grandChild1"] = grandChild1
child2.Groups["grandChild2"] = grandChild2
child2.Groups["grandChild3"] = grandChild3
// The resulting config structure of groups looks like:
// root:
       child1:
//
            grandChild1
//
//
      child2:
        grandChild2
//
            grandChild3
//
```

Each group defines a level in the config hierarchy, and each group has an associated set of values (indexed by string key) and policies (also indexed by string key).

Values are defined by:

```
message ConfigValue {
    uint64 version = 1;
    bytes value = 2;
    string mod_policy = 3;
}
```

```
message ConfigPolicy {
    uint64 version = 1;
    Policy policy = 2;
    string mod_policy = 3;
}
```

Note that Values, Policies, and Groups all have a version and a mod_policy. The version of an element is incremented each time that element is modified. The mod_policy is used to govern the required signatures to modify that element. For Groups, modification is adding or removing elements to the Values, Policies, or Groups maps (or changing the mod_policy). For Values and Policies, modification is changing the Value and Policy fields respectively (or changing the mod_policy). Each element's mod_policy is evaluated in the context of the current level of the config. Consider the following example mod policies defined at Channel.Groups["Application"] (Here, we use the golang map reference syntax, so Channel.Groups["Application"].Policies["policy1"] refers to the base Channel group's Application group's Policies map's policy1 policy.)

```
• policy1 maps to Channel.Groups["Application"].Policies["policy1"]
```

- Org1/policy2 maps to Channel.Groups["Application"].Groups["Org1"].Policies["policy2"]
- /Channel/policy3 maps to Channel.Policies["policy3"]

Note that if a mod_policy references a policy which does not exist, the item cannot be modified.

Configuration updates

```
Configuration updates are submitted as an Envelope message of type

HeaderType_CONFIG_UPDATE. The Payload data of the transaction is a marshaled

ConfigUpdateEnvelope. The ConfigUpdateEnvelope is defined as follows:
```

```
message ConfigUpdateEnvelope {
   bytes config_update = 1;
   repeated ConfigSignature signatures = 2;
}
```

The signatures field contains the set of signatures which authorizes the config update. Its message definition is:

```
message ConfigSignature {
   bytes signature_header = 1;
   bytes signature = 2;
}
```

The signature_header is as defined for standard transactions, while the signature is over the concatenation of the signature_header bytes and the config_update bytes from the ConfigUpdateEnvelope message.

The ConfigUpdateEnvelope config_update bytes are a marshaled ConfigUpdate message which is defined as follows:

```
message ConfigUpdate {
    string channel_id = 1;
    ConfigGroup read_set = 2;
    ConfigGroup write_set = 3;
}
```

The channel_id is the channel ID the update is bound for, this is necessary to scope the signatures which support this reconfiguration.

The read_set specifies a subset of the existing configuration, specified sparsely where only the version field is set and no other fields must be populated. The particular configValue value or configPolicy policy fields should never be set in the read_set. The configGroup may have a subset of its map fields populated, so as to reference an element deeper in the config tree. For instance, to include the Application group in the read_set, its parent (the channel group) must also be included in the read set, but, the channel group does not need to populate all of the keys, such as the orderer group key, or any of the values or policies keys.

The write_set specifies the pieces of configuration which are modified. Because of the hierarchical nature of the configuration, a write to an element deep in the hierarchy must contain the higher level elements in its write_set as well. However, for any element in the write_set which is also specified in the read_set at the same version, the element should be specified sparsely, just as in the read_set.

For example, given the configuration:

```
Channel: (version 0)
Orderer (version 0)
Application (version 3)
Org1 (version 2)
```

To submit a configuration update which modifies <code>Org1</code> , the <code>read_set</code> would be:

```
Channel: (version 0)
Application: (version 3)
```

and the write_set would be

```
Channel: (version 0)
Application: (version 3)
Org1 (version 3)
```

When the **CONFIG_UPDATE** is received, the orderer computes the resulting **CONFIG** by doing the following:

- 1. Verifies the channel_id and read_set must exist at the given versions.
- 2. Computes the update set by collecting all elements in the write_set which do not appear at the same version in the read_set.
- 3. Verifies that each element in the update set increments the version number of the element update by exactly 1.
- 4. Verifies that the signature set attached to the ConfigUpdateEnvelope satisfies the mod policy for each element in the update set.
- 5. Computes a new complete version of the config by applying the update set to the current config.
- 6. Writes the new config into a ConfigEnvelope which includes the CONFIG_UPDATE as the last_update field and the new config encoded in the config field, along with the incremented sequence value.
- 7. Writes the new ConfigEnvelope into a Envelope of type CONFIG, and ultimately writes this as the sole transaction in a new configuration block.

When the peer (or any other receiver for <code>Deliver</code>) receives this configuration block, it should verify that the config was appropriately validated by applying the <code>last_update</code> message to the current config and verifying that the orderer-computed <code>config</code> field contains the correct new configuration.

Permitted configuration groups and values

Any valid configuration is a subset of the following configuration. Here we use the notation <code>peer.<MSG></code> to define a <code>ConfigValue</code> whose <code>value</code> field is a marshaled proto message of name <code><MSG></code> defined in <code>fabric/protos/peer/configuration.proto</code>. The notations <code>common.<MSG></code>, <code>msp.<MSG></code>, and <code>orderer.<MSG></code> correspond similarly, but with their messages defined in <code>fabric/protos/common/configuration.proto</code>, <code>fabric/protos/msp/mspconfig.proto</code>, and <code>fabric/protos/orderer/configuration.proto</code> respectively.

Note, that the keys {{org_name}} and {{consortium_name}} represent arbitrary names, and indicate an element which may be repeated with different names.

```
&ConfigGroup{
    Groups: map<string, *ConfigGroup> {
        "Application": &ConfigGroup{
            Groups:map<String, *ConfigGroup> {
                {{org_name}}:&ConfigGroup{
                    Values:map<string, *ConfigValue>{
                        "MSP":msp.MSPConfig,
                        "AnchorPeers":peer.AnchorPeers,
                    },
                },
            },
        "Orderer":&ConfigGroup{
            Groups:map<String, *ConfigGroup> {
                {{org_name}}:&ConfigGroup{
                    Values:map<string, *ConfigValue>{
                        "MSP":msp.MSPConfig,
                    },
                },
            },
            Values:map<string, *ConfigValue> {
                "ConsensusType":orderer.ConsensusType,
                "BatchSize": orderer. BatchSize,
                "BatchTimeout": orderer.BatchTimeout,
                "KafkaBrokers":orderer.KafkaBrokers,
            },
        Groups:map<String, *ConfigGroup> {
                {{consortium_name}}:&ConfigGroup{
                    Groups:map<string, *ConfigGroup> {
                        {{org name}}:&ConfigGroup{
                            Values:map<string, *ConfigValue>{
                                "MSP":msp.MSPConfig,
                            },
                        },
                    },
                    Values:map<string, *ConfigValue> {
                        "ChannelCreationPolicy":common.Policy,
               },
           },
       },
   Values: map<string, *ConfigValue> {
        "HashingAlgorithm":common.HashingAlgorithm,
        "BlockHashingDataStructure":common.BlockDataHashingStructure,
        "Consortium":common.Consortium,
        "OrdererAddresses":common.OrdererAddresses,
   },
}
```

Orderer system channel configuration

The ordering system channel needs to define ordering parameters, and consortiums for creating channels. There must be exactly one ordering system channel for an ordering service, and it is the first channel to be created (or more accurately bootstrapped). It is recommended never to define an Application section inside of the ordering system channel genesis configuration, but may be done for testing. Note that any member with read access to the ordering system channel may see all channel creations, so this channel's access should be restricted.

The ordering parameters are defined as the following subset of config:

```
&ConfigGroup{
    Groups: map<string, *ConfigGroup> {
        "Orderer": &ConfigGroup{
            Groups:map<String, *ConfigGroup> {
                {{org_name}}:&ConfigGroup{
                    Values:map<string, *ConfigValue>{
                         "MSP":msp.MSPConfig,
                    },
                },
            },
            Values:map<string, *ConfigValue> {
                "ConsensusType":orderer.ConsensusType,
                "BatchSize":orderer.BatchSize,
                "BatchTimeout":orderer.BatchTimeout,
                "KafkaBrokers": orderer. KafkaBrokers,
            },
        },
    },
```

Each organization participating in ordering has a group element under the Orderer group. This group defines a single parameter MSP which contains the cryptographic identity information for that organization. The Values of the Orderer group determine how the ordering nodes function. They exist per channel, so Orderer.BatchTimeout for instance may be specified differently on one channel than another.

At startup, the orderer is faced with a filesystem which contains information for many channels. The orderer identifies the system channel by identifying the channel with the consortiums group defined. The consortiums group has the following structure.

```
&ConfigGroup{
    Groups: map<string, *ConfigGroup> {
        "Consortiums": &ConfigGroup{
            Groups:map<String, *ConfigGroup> {
                {{consortium name}}:&ConfigGroup{
                    Groups:map<string, *ConfigGroup> {
                        {{org_name}}:&ConfigGroup{
                            Values:map<string, *ConfigValue>{
                                "MSP":msp.MSPConfig,
                            },
                        },
                    Values:map<string, *ConfigValue> {
                        "ChannelCreationPolicy":common.Policy,
               },
      },
   },
},
```

Note that each consortium defines a set of members, just like the organizational members for the ordering orgs. Each consortium also defines a ChannelCreationPolicy. This is a policy which is applied to authorize channel creation requests. Typically, this value will be set to an ImplicitMetaPolicy requiring that the new members of the channel sign to authorize the channel creation. More details about channel creation follow later in this document.

Application channel configuration

Application configuration is for channels which are designed for application type transactions. It is defined as follows:

Just like with the Orderer section, each organization is encoded as a group. However, instead of only encoding the MSP identity information, each orgadditionally encodes a list of AnchorPeers. This list allows the peers of different organizations to contact each other for peer gossip networking.

The application channel encodes a copy of the orderer orgs and consensus options to allow for deterministic updating of these parameters, so the same <u>Orderer</u> section from the orderer system channel configuration is included. However from an application perspective this may be largely ignored.

Channel creation

When the orderer receives a **CONFIG_UPDATE** for a channel which does not exist, the orderer assumes that this must be a channel creation request and performs the following.

- 1. The orderer identifies the consortium which the channel creation request is to be performed for. It does this by looking at the **Consortium** value of the top level group.
- 2. The orderer verifies that the organizations included in the Application group are a subset of the organizations included in the corresponding consortium and that the ApplicationGroup is set to version 1.
- 3. The orderer verifies that if the consortium has members, that the new channel also has application members (creation consortiums and channels with no members is useful for testing only).
- 4. The orderer creates a template configuration by taking the Orderer group from the ordering system channel, and creating an Application group with the newly specified members and specifying its mod_policy to be the ChannelCreationPolicy as specified in the consortium config. Note that the policy is evaluated in the context of the new

- configuration, so a policy requiring ALL members, would require signatures from all the new channel members, not all the members of the consortium.
- 5. The orderer then applies the CONFIG_UPDATE as an update to this template configuration. Because the CONFIG_UPDATE applies modifications to the Application group (its version is 1), the config code validates these updates against the ChannelCreationPolicy. If the channel creation contains any other modifications, such as to an individual org's anchor peers, the corresponding mod policy for the element will be invoked.
- 6. The new **CONFIG** transaction with the new channel config is wrapped and sent for ordering on the ordering system channel. After ordering, the channel is created.